

# CNN and GAN with GPU Acceleration - Final Report

## Convolutional Neural Networks (CNNs) with GPU Acceleration

### Overview:

-----

CNNs are particularly well-suited for image processing tasks due to their ability to capture spatial hierarchies in data.

When trained using GPUs, CNNs benefit immensely from parallelization, especially in convolutional and matrix operations.

### Experiment Setup:

-----

- Dataset: Cats vs Dogs (Kaggle)
- Input: 150x150 grayscale images
- Architecture:

Conv2D -> ReLU -> MaxPooling -> Conv2D -> ReLU -> MaxPooling -> Flatten -> Dense -> Output

- Activation: ReLU, Sigmoid
- Loss: Binary Crossentropy
- Optimizer: Adam

### GPU Utilization:

-----

Convolution and matrix multiplication layers utilize the GPU cores in parallel. This leads to much faster training

compared to CPUs, particularly as batch sizes and model complexity increase.

### Visualizing Filters:

-----  
Using Keras, we accessed and visualized the first layer's learned filters:

```
filters, _ = model.layers[0].get_weights()
```

This allowed us to inspect what the model was learning in early convolutional layers.

Results:

- 
- Epochs: 10
  - Accuracy: ~85% with minimal tuning
  - Training Time: ~5x faster on GPU vs CPU

Conclusion:

-----  
CNNs are one of the best examples of GPU utilization in AI, due to highly parallelizable operations.  
Generative Adversarial Networks (GANs) with GPU Acceleration

Overview:

-----  
GANs consist of two models: a Generator (G) and a Discriminator (D), trained in opposition.

The Generator learns to create realistic data while the Discriminator learns to distinguish fake from real.

Why GPU Acceleration?

- 
- GANs train slowly due to multiple forward and backward passes for both G and D.
  - Parallel GPU cores help speed up training of both networks simultaneously.

## Experiment Setup:

-----

- Dataset: MNIST (handwritten digits)
- Generator: Dense layers with ReLU/Tanh
- Discriminator: Dense layers with LeakyReLU/Sigmoid
- Loss: Binary Crossentropy
- Optimizer: Adam

## Training Loop:

-----

Each training step involves:

1. Generating fake images from random noise (G)
2. Training D on real and fake images
3. Training G to fool D

# Sample generator input

```
noise = np.random.normal(0, 1, (batch_size, latent_dim))
```

```
generated_images = generator.predict(noise)
```

## Results:

-----

- Trained for 10,000 steps
- Generated realistic MNIST digits
- Training Time: ~4x faster on GPU

## Conclusion:

-----

GANs demand significant compute, making them a perfect use-case for GPU acceleration.